

```

# Importing the Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

bmi = pd.read_csv(r"C:\Users\Neil\Supervised Machine Learning Lab
(SMLL)\BMI.csv")
bmi.head()

```

	Gender	Height	Weight	Final_Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3

```

bmi.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 4 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Gender          500 non-null   object
 1   Height          500 non-null   int64
 2   Weight          500 non-null   int64
 3   Final_Index     500 non-null   int64
dtypes: int64(3), object(1)
memory usage: 15.8+ KB

bmi.shape

(500, 4)

bmi.isna().sum()

Gender          0
Height          0
Weight          0
Final_Index     0
dtype: int64

# Create a figure with 3 subplots
fig, ax = plt.subplots(1, 3, figsize=(15, 5))

# Plot histogram for 'Height'
ax[0].hist(bmi['Height'], bins=10, color='skyblue', edgecolor='black')
ax[0].set_title('Height Distribution')
ax[0].set_xlabel('Height')
ax[0].set_ylabel('Frequency')

```


```

# Plot histogram for 'Weight'
ax[1].hist(bmi['Weight'], bins=10, color='lightgreen',
edgecolor='black')
ax[1].set_title('Weight Distribution')
ax[1].set_xlabel('Weight')
ax[1].set_ylabel('Frequency')

# Plot histogram for 'Final_Index'
ax[2].hist(bmi['Final_Index'], bins=10, color='salmon',
edgecolor='black')
ax[2].set_title('Final Index Distribution')
ax[2].set_xlabel('Final Index')
ax[2].set_ylabel('Frequency')

# Adjust layout and show the plot
plt.tight_layout()
plt.show()

```



```

# Map the Gender column
bmi['Gender'] = bmi['Gender'].map({'Male': 0, 'Female': 1})

# Display the updated DataFrame
print(bmi.head())

```

	Gender	Height	Weight	Final_Index
0	0	174	96	4
1	0	189	87	2
2	1	185	110	4
3	1	195	104	3
4	0	149	61	3

```

# Vif dataframe
vif_data = pd.DataFrame()
vif_data['Feature'] = ['Gender', 'Height', 'Weight']

vif_data.head()

```

	Feature
0	Gender
1	Height
2	Weight

```
from statsmodels.stats.outliers_influence import
variance_inflation_factor
vif_data['VIF'] = [variance_inflation_factor(bmi.values,i) for i in
range(3)]
```

```
# Display the VIF DataFrame
print(vif_data)
```

	Feature	VIF
0	Gender	2.031010
1	Height	11.941301
2	Weight	40.024880

```
bmi['Height_m'] = bmi['Height'] * 0.01
```

```
y = bmi['Final_Index']
```

```
print(bmi.head())
```

	Gender	Height	Weight	Final_Index	Height_m
0	0	174	96	4	1.74
1	0	189	87	2	1.89
2	1	185	110	4	1.85
3	1	195	104	3	1.95
4	0	149	61	3	1.49

```
bmi.drop(columns=['Height'], inplace=True)
```

```
bmi.head()
```

	Gender	Weight	Final_Index	Height_m
0	0	96	4	1.74
1	0	87	2	1.89
2	1	110	4	1.85
3	1	104	3	1.95
4	0	61	3	1.49

```
bmi.drop(columns=['Final_Index'], inplace=True)
```

```
bmi
```

	Gender	Weight	Height_m
0	0	96	1.74
1	0	87	1.89
2	1	110	1.85
3	1	104	1.95
4	0	61	1.49

..
495	1	153	1.50
496	1	121	1.84
497	1	136	1.41
498	0	95	1.50
499	0	131	1.73

[500 rows x 3 columns]

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
X = bmi
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# importing the model
mlr = LinearRegression()
# Fitting the model
mlr.fit(X_train,y_train)
```

```
LinearRegression()
```

```
# Predicting
y_pred = mlr.predict(X_test)
print(y_pred)
```

```
[4.00042305 2.35182455 3.25757795 3.28935721 4.09980251 6.01906252
 4.73630207 5.63107258 3.10299365 5.82847118 3.64836074 4.40777202
 1.89479948 4.43947187 4.69633276 2.80201508 6.05634566 4.90894203
 2.89304327 5.16442014 2.30908967 5.99284168 1.53994432 3.65522254
 3.8913392 5.32453559 5.79265024 6.26204098 4.32076064 2.45683944
 4.1771071 5.44603967 3.98938802 3.64554302 3.31557805 2.12543954
 4.79158873 2.42785303 3.15963582 2.26896154 4.76946412 1.9762797
 3.84575933 3.06291766 4.93795572 3.23553275 3.29209551 3.21346028
 3.47577292 3.44125535 3.92452612 3.75741731 1.86307477 1.73189117
 5.79126745 2.51890849 2.60722324 6.06046675 6.54232986 4.052892
 4.1509384 5.353522 4.34416377 3.91761218 1.48750265 2.31456868
 3.09193134 3.62208777 2.15719153 2.11437723 2.17655056 4.0943235
 2.16540884 3.17891544 5.78709181 4.9089693 5.30241098 3.88993155
 2.56446108 4.93654565 4.42710619 2.96348846 2.43883592 5.44186403
 3.49369702 5.08565335 2.89865384 4.07916497 4.40766774 2.63902737
 1.50819233 2.01359012 6.19303313 3.26871967 5.62415864 2.4112323
 4.74462607 4.40354665 1.71944608 4.36083664]
```

```
# Actual vs Predicted value
mlr_diff = pd.DataFrame({ "Actual Value " : y_test, "Predicted value "
: y_pred})
print(mlr_diff)
```

	Actual Value	Predicted value
361	4	4.000423
73	2	2.351825
374	4	3.257578
155	4	3.289357
104	4	4.099803
..
347	2	2.411232
86	5	4.744626
75	4	4.403547
438	1	1.719446
15	5	4.360837

[100 rows x 2 columns]

Errors

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

Mean Square Error

```
mean_square_error = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Square Error is : {mean_square_error}")
```

Calculate Mean Absolute Error

```
mae = mean_absolute_error(y_test, y_pred)
```

```
print(f"Mean Absolute Error is : {mae}")
```

Root Mean Square Error

```
root_mean_square_error = np.sqrt(mean_square_error)
```

```
print(f"Root Mean Square Error is : {root_mean_square_error}")
```

Calculate R-squared

```
r2 = r2_score(y_test, y_pred)
```

```
print(f"R2 score is : {r2*100}")
```

#intecept

```
X_intercept = mlr.intercept_
```

```
print(f"X intercept is : {X_intercept}")
```

```
Y_intercept = mlr.coef_
```

```
print(f"Y intercept is : {Y_intercept}")
```

Mean Square Error is : 0.3393556144069992

Mean Absolute Error is : 0.4684465810376227

Root Mean Square Error is : 0.5825423713404881

R2 score is : 79.62929261018073

X intercept is : 6.225070667746884

Y intercept is : [-0.06485949 0.03451756 -3.59003525]

Function to scatter plot with regression line

```
def
```

```
plot_scatter_with_regression(ax, x_data, y_data, x_label, y_label, title):
    ax.scatter(x_data, y_data, alpha=0.5)
```

Calculate regression line

```
slope, intercept = np.polyfit(x_data, y_data, 1)
```

```

    regression_line = slope*x_data + intercept

    ax.plot(x_data, regression_line, color='red',
            label='Regression_Line')

    ax.set_title(title)
    ax.set_xlabel(x_label)
    ax.set_ylabel(y_label)
    ax.legend()

bmi.columns

Index(['Gender', 'Weight', 'Height_m'], dtype='object')

fig, ax = plt.subplots(1, 3, figsize=(20, 10))

# Scatter plot for Gender vs Final_Index
plot_scatter_with_regression(ax[0], bmi['Gender'], y, 'Gender', 'Final
Index', 'Gender vs Final Index')

# Scatter plot for Weight vs Final_Index
plot_scatter_with_regression(ax[1], bmi['Weight'], y, 'Weight', 'Final
Index', 'Weight vs Final Index')

# Scatter plot for Height vs Final_Index
plot_scatter_with_regression(ax[2], bmi['Height_m'], y, 'Height',
'Final Index', 'Height vs Final Index')

plt.tight_layout()
plt.show()

```

