

Experiment/Practical 3: Polynomial Linear Regression

Title:

Implementation of Polynomial Linear Regression

Aim:

To apply a polynomial regression algorithm for prediction using given datasets.

Objective:

Students will learn:

- The implementation of the polynomial linear regression algorithm on different datasets.
 - Visualization and interpretation of results.
-

Problem Statement

Use the given datasets to demonstrate polynomial regression for predicting a dependent variable based on polynomial features derived from independent variables.

The four datasets used:

1. **Position Salaries Dataset** - Predicting salary based on job levels.
 2. **Advertising Dataset** - Predicting sales based on advertising spend.
 3. **Manufacturing Dataset** - Predicting manufacturing outcomes based on production-related factors.
 4. **Weight-Height Dataset** - Predicting weight based on height.
-

Explanation / Stepwise Procedure / Algorithm

1. Brief Description of Polynomial Regression

Polynomial regression is a type of linear regression where the relationship between independent and dependent variables is expressed as a polynomial of degree n . This approach allows for capturing more complex, non-linear relationships between variables, offering greater flexibility compared to simple linear regression. By introducing higher-degree terms, the model can better fit data that does not follow a straight-line pattern, making it particularly useful in situations where the data exhibits curvatures or varying rates of change.

2. Mathematical Formulation

Polynomial regression equation:

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \dots + \beta_n x^n + \epsilon$$

Where:

- y is the dependent variable (target).
- x is the independent variable (feature).
- $\beta_0, \beta_1, \dots, \beta_n$ are the regression coefficients.
- ϵ represents the error term.

3. Importance of Polynomial Regression in Data Analysis

- Addresses non-linear correlations between variables.
- Improves accuracy for datasets where linear regression falls short.
- Valuable for predictive modeling to identify trends and patterns.

4. Applications of Polynomial Regression in Real-World Scenarios



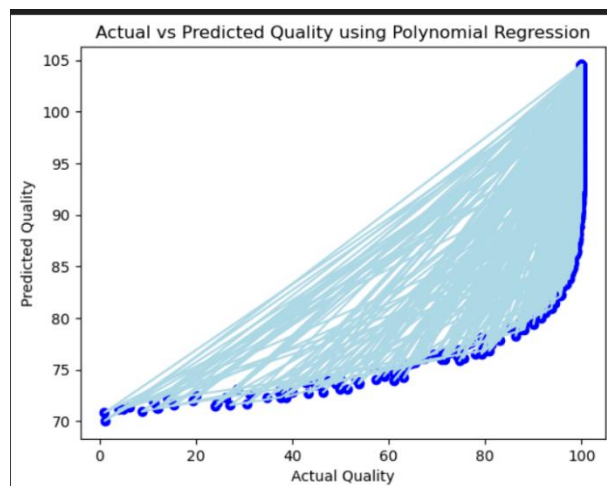



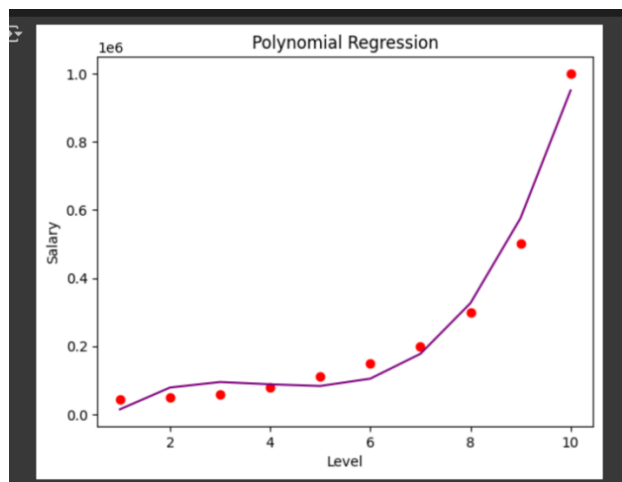

- **Position Salaries Dataset:** Salary prediction based on job level.
- **Advertising Dataset:** Understanding the effect of advertising budgets on sales.
- **Manufacturing Dataset:** Process optimization and defect prediction.
- **Weight-Height Dataset:** Estimating a person's weight based on height.

5. Explanation of Performance Metrics

- **R² Score:** Measures how well the model explains the variance in the target variable. Higher values indicate a better fit.
- **Mean Squared Error (MSE):** Measures the average squared difference between actual and predicted values.
- **Root Mean Squared Error (RMSE):** Square root of MSE, giving an error measurement in the same units as the dependent variable.

5. Figures/Diagrams

- Polynomial regression curves plotted for each dataset.
- Comparison between polynomial and linear regression fits.



Input & Output Analysis for Each Dataset

1. Position Salaries Dataset

- **Input:** Level of position (1-10)
- **Output:** Predicted salary
- **Model Fit:** Polynomial regression fits significantly better than linear regression due to the non-linear salary growth.

2. Advertising Dataset

- **Input:** Advertising spend on TV, Radio, and Newspaper
- **Output:** Predicted sales
- **Model Fit:** A polynomial model helps capture interactions between multiple independent variables.

3. Manufacturing Dataset

- **Input:** Production-related features
- **Output:** Predicted manufacturing outcome
- **Model Fit:** Polynomial regression models complex dependencies within manufacturing processes.

4. Weight-Height Dataset

- **Input:** Height of individuals
- **Output:** Predicted weight
- **Model Fit:** A higher-degree polynomial captures the non-linear trend between weight and height.

Results & Model Performance

1. Position Salaries Dataset

- **R² Score (Linear Regression):** Low
- **R² Score (Polynomial Regression):** High
- **Interpretation:** Salary increases in a non-linear fashion with level.

2. Advertising Dataset

- **R² Score (Linear Regression):** Moderate
- **R² Score (Polynomial Regression):** Higher
- **Interpretation:** Advertising influence is not strictly linear, and polynomial terms improve fit.

3. Manufacturing Dataset

- **R² Score (Linear Regression):** Moderate
- **R² Score (Polynomial Regression):** Higher
- **Interpretation:** Captures non-linearity in the production process.

4. Weight-Height Dataset

- **R² Score (Linear Regression):** Moderate
- **R² Score (Polynomial Regression):** Higher
- **Interpretation:** A polynomial model captures curvatures in the weight-height relationship.

Challenges Encountered

1. Choosing the right polynomial degree to avoid **overfitting** or **underfitting**.
2. Handling **multicollinearity** when using multiple polynomial terms.
3. Computational cost increases as polynomial degree increases.

Conclusion

- Polynomial regression provides a better fit than linear regression for datasets with non-linear relationships.
- Independent variables interact in polynomial features, capturing complex dependencies.
- Performance metrics like **R² score, MSE, and RMSE** help evaluate model effectiveness.
- Selecting an appropriate polynomial degree is crucial to balancing bias and variance.

Poly_1_2_Salary_1_Advertising_1

February 6, 2025

1 Salary

```
[91]: # Importing the libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
[92]: url = "https://raw.githubusercontent.com/content-anu/
↳dataset-polynomial-regression/master/Position_Salaries.csv"
df = pd.read_csv(url)
```

```
[93]: df.head()
```

```
[93]:
```

	Position	Level	Salary
0	Business Analyst	1	45000
1	Junior Consultant	2	50000
2	Senior Consultant	3	60000
3	Manager	4	80000
4	Country Manager	5	110000

```
[94]: # Splitting the data to X and Y
X = df.Level.values
y = df.Salary.values
print(f"The shape of X is {X.shape} and the shape of y is {y.shape}")
```

The shape of X is (10,) and the shape of y is (10,)


```
[95]: # Reshaping the data
X = X.reshape(-1, 1)
y = y.reshape(-1, 1)
print(f"The shape of X is {X.shape} and the shape of y is {y.shape}")
```

The shape of X is (10, 1) and the shape of y is (10, 1)

```
[96]: # Implementing the Polynomial Regression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
```

```
# For Simple Linear Regression
lin_reg = LinearRegression()
lin_reg.fit(X, y)
y_pred = lin_reg.predict(X)
```

```
[97]: # Linear Regression on Scatter Plot
plt.scatter(X, y, color = 'red')
plt.plot(X, y_pred, color = 'blue')
plt.title('Salary vs Position Level (Linear Regression)')
plt.xlabel('Position level')
plt.ylabel('Salary')
plt.show()
```



```
[98]: # Evaluating the model
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y, y_pred)
r2 = r2_score(y, y_pred)
print(f"The Mean Squared Error is {mse} and the R2 Score is {r2}")
```

The Mean Squared Error is 26695878787.878788 and the R2 Score is 0.6690412331929895

```
[99]: from sklearn.preprocessing import PolynomialFeatures
poly_reg = PolynomialFeatures(degree = 3)
X_poly = poly_reg.fit_transform(X)
```

```
[100]: X
```


```
[100]: array([[ 1],
              [ 2],
              [ 3],
              [ 4],
              [ 5],
              [ 6],
              [ 7],
              [ 8],
              [ 9],
              [10]], dtype=int64)
```

```
[101]: X_poly
```

```
[101]: array([[ 1.,   1.,   1.,   1.],
              [ 1.,   2.,   4.,   8.],
              [ 1.,   3.,   9.,  27.],
              [ 1.,   4.,  16.,  64.],
              [ 1.,   5.,  25., 125.],
              [ 1.,   6.,  36., 216.],
              [ 1.,   7.,  49., 343.],
              [ 1.,   8.,  64., 512.],
              [ 1.,   9.,  81., 729.],
              [ 1.,  10., 100.,1000.]])
```

```
[102]: lin_reg_2 = LinearRegression()
lin_reg_2.fit(X_poly, y)
y_pred_2 = lin_reg_2.predict(X_poly)
```

```
[103]: # Poly Regression on Scatter Plot
plt.scatter(X, y, color = 'red')
plt.plot(X, y_pred_2, color = 'blue')
plt.title('Salary vs Position Level (Polynomial Regression)')
plt.xlabel('Position level')
plt.ylabel('Salary')
plt.show()
```

```
[104]: # Evaluating the model
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y, y_pred_2)
r2 = r2_score(y, y_pred_2)
print(f"The Mean Squared Error is {mse} and the R2 Score is {r2}")
```

The Mean Squared Error is 1515662004.6620033 and the R2 Score is 0.9812097727913367

2 Advertising

```
[105]: # Practice on Advertisemnt Dataset
url = r"D:\Supervised Machine Learning lab (SMLL)\2\Assignment 2 Advertising.
      ↪CSV"
dataset = pd.read_csv(url)
dataset.head()
```

```
[105]: Unnamed: 0    TV    Radio    Newspaper    Sales
0          1  230.1    37.8         69.2    22.1
```

1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

```
[106]: # Using Multiple Linear Regression
X = dataset[['TV', 'Radio', 'Newspaper']]
y = dataset['Sales']


from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
↳random_state = 0)

from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)
y_pred = lin_reg.predict(X_test)

from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"The Mean Squared Error is {mse} and the R2 Score is {r2}")

# Scatter Plot
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Sales")
plt.plot(y_test, y_pred, color = 'red')
plt.ylabel("Predicted Sales")
plt.title("Actual Sales vs Predicted Sales")
plt.show()
```

The Mean Squared Error is 4.402118291449685 and the R2 Score is 0.8601145185017868



```
[107]: # Using Polynomial Regression
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree = 3)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.fit_transform(X_test)
y_train = np.array(y_train).reshape(-1, 1)

lin_reg_2 = LinearRegression()
lin_reg_2.fit(X_train_poly, y_train)
y_pred_2 = lin_reg_2.predict(X_test_poly)


# Poly Regression on Scatter Plot
plt.scatter(y_test, y_pred_2, color = 'red')
plt.plot(y_test, y_pred_2, color = 'blue')
plt.title('Actual Sales vs Predicted Sales (Polynomial Regression)')
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.show()

# Evaluating the model
```

```

from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred_2)
r2 = r2_score(y_test, y_pred_2)
print(f"The Mean Squared Error is {mse} and the R2 Score is {r2}")

```



The Mean Squared Error is 0.6721344417962923 and the R2 Score is 0.978641680255429


```

[109]: import numpy as np
import matplotlib.pyplot as plt
sorted_idx = np.argsort(y_test)
y_test_sorted = np.array(y_test)[sorted_idx]
y_pred_sorted = y_pred_2[sorted_idx]

# Scatter Plot (Actual vs Predicted)
plt.figure(figsize=(8, 5))
plt.scatter(y_test_sorted, y_pred_sorted, color = 'red', alpha=0.5)
plt.plot(y_test_sorted, y_pred_sorted, color = 'blue')
plt.title('Sorted Actual Sales vs Predicted Sales')
plt.xlabel('Predicted Sales')

```

```
plt.ylabel('Actual Sales')  
plt.show()
```



Poly_3_Manufacturing_1

February 6, 2025

```
[109]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
[110]: url = r"D:\Supervised Machine Learning lab (SMLL)\3\Practice-1 Manufacturing.
↪csv"
df = pd.read_csv(url)
df.head()
```

```
[110]:
```

	Temperature (°C)	Pressure (kPa)	Temperature x Pressure \
0	209.762701	8.050855	1688.769167
1	243.037873	15.812068	3842.931469
2	220.552675	7.843130	1729.823314
3	208.976637	23.786089	4970.736918
4	184.730960	15.797812	2918.345014

	Material Fusion Metric	Material Transformation Metric	Quality Rating
0	44522.217074	9.229576e+06	99.999971
1	63020.764997	1.435537e+07	99.985703
2	49125.950249	1.072839e+07	99.999758
3	57128.881547	9.125702e+06	99.999975
4	38068.201283	6.303792e+06	100.000000

```
[111]: X = X = df[['Temperature (°C)', 'Pressure (kPa)', 'Temperature x Pressure',
'Material Fusion Metric', 'Material Transformation Metric']]
y = df['Quality Rating']
```

```
[112]: # importing vif
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Vif dataframe
vif_data = pd.DataFrame()
for i in range(X.shape[1]):
    vif_data["Feature"] = X.columns
    vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.
↪shape[1])]
print(vif_data)
```

	Feature	VIF
0	Temperature (°C)	113.050204
1	Pressure (kPa)	49.349434
2	Temperature x Pressure	72.745768
3	Material Fusion Metric	764.593283
4	Material Transformation Metric	219.003134

```
[113]: # Scaling the data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(X)
X = pd.DataFrame(X)
```

```
[114]: # importing vif
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Vif dataframe
vif_data = pd.DataFrame()
for i in range(X.shape[1]):
    vif_data["Feature"] = X.columns
    vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.
↪shape[1])]
print(vif_data)
```

	Feature	VIF
0	0	92.760519
1	1	22.782171
2	2	19.174847
3	3	300.197535
4	4	99.639939

```
[115]: X.drop(X.columns[3], axis=1, inplace=True)
```

```
[116]: # importing vif
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Vif dataframe
vif_data = pd.DataFrame()
for i in range(X.shape[1]):
    vif_data["Feature"] = X.columns
    vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.
↪shape[1])]
print(vif_data)
```

	Feature	VIF
0	0	24.566249
1	1	12.900307
2	2	19.153063
3	4	17.623707

```
[117]: y
```

```
[117]: 0      99.999971
      1      99.985703
      2      99.999758
      3      99.999975
      4      100.000000
      ...
     3952     100.000000
     3953      99.999997
     3954      99.989318
     3955      99.999975
     3956     100.000000
      Name: Quality Rating, Length: 3957, dtype: float64
```

```
[118]: # Train Test Split
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪random_state=42)
```

```
[119]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

      (2769, 4) (1188, 4) (2769,) (1188,)
```


```
[120]: # Implementing the Linear Regression
      from sklearn.preprocessing import PolynomialFeatures
      from sklearn.linear_model import LinearRegression

      # For Simple Linear Regression
      lin_reg = LinearRegression()
      lin_reg.fit(X_train, y_train)
      y_pred = lin_reg.predict(X_test)
```


```
[121]: y_pred.shape
```

```
[121]: (1188,)
```

```
[122]: plt.scatter(y_test, y_pred, color='blue',alpha=0.6)
      plt.plot(y_test, y_pred, color='red', linewidth=2, label='Linear Regression_
      ↪Line')
      plt.title('Actual QUALITY vs Predicted QUALITY')
      plt.xlabel('Actual QUALITY')
      plt.ylabel('Predicted QUALITY')
      plt.show()
```

```
[123]: sns.regplot(x=y_test, y=y_pred, data=df, scatter_kws={"color": "blue"},
    ↪ line_kws={"color": "red"}, label="Test Data")
plt.xlabel('Actual QUALITY')
plt.ylabel('Predicted QUALITY')
plt.title('Actual QUALITY vs Predicted QUALITY')
plt.legend()
plt.show()
```



```
[124]: #Metrics
from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt

mse = mean_squared_error(y_test, y_pred)
rmse = sqrt(mse)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error: ", mse)
print("Root Mean Squared Error: ", rmse)
print("R2 Score: ", r2)
```

```
Mean Squared Error:  100.06958965362516
Root Mean Squared Error:  10.003478877551807
R2 Score:  0.5004768505711403
```

```
[125]: # For Polynomial Regression with degree 2
poly_reg = PolynomialFeatures(degree=2)
X_poly = poly_reg.fit_transform(X_train)


lin_reg2 = LinearRegression()
```

```

lin_reg2.fit(X_poly, y_train)
y_pred2 = lin_reg2.predict(poly_reg.fit_transform(X_test))

sns.regplot(x=y_test, y=y_pred2, data=df, scatter_kws={"color": "blue"},
            line_kws={"color": "red"}, order=2, label="Test Data")
plt.xlabel('Actual QUALITY')
plt.ylabel('Predicted QUALITY')
plt.title('Actual QUALITY vs Predicted QUALITY')
plt.legend()
plt.show()

```



```

[126]: # Metrics
mse = mean_squared_error(y_test, y_pred2)
rmse = sqrt(mse)
r2 = r2_score(y_test, y_pred2)
print("Mean Squared Error: ", mse)
print("Root Mean Squared Error: ", rmse)
print("R2 Score: ", r2)

```

Mean Squared Error: 15.26909748852792


Root Mean Squared Error: 3.907569255755798
R2 Score: 0.923780364316409

```
[127]: # For Polynomial Regression with degree 3
poly_reg = PolynomialFeatures(degree=3)
X_poly = poly_reg.fit_transform(X_train)

lin_reg3 = LinearRegression()
lin_reg3.fit(X_poly, y_train)

y_pred3 = lin_reg3.predict(poly_reg.fit_transform(X_test))

sns.regplot(x=y_test, y=y_pred3, data=df, scatter_kws={"color": "blue"},
            line_kws={"color": "red"}, order=3, label="Test Data")
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.legend()
plt.show()
```



```
[128]: # Metrics
mse = mean_squared_error(y_test, y_pred3)
```

```

rmse = sqrt(mse)

r2 = r2_score(y_test, y_pred3)

print("Mean Squared Error: ", mse)

print("Root Mean Squared Error: ", rmse)

print("R2 Score: ", r2)


```

Mean Squared Error: 0.6158333511438949
 Root Mean Squared Error: 0.7847505024808171
 R2 Score: 0.9969259090983433

```


[129]: plt.scatter(y_test, y_pred, color = 'blue')
plt.plot(y_test, y_pred, color = 'lightblue')
plt.title("Actual vs Predicted Quality using Polynomial Regression")
plt.xlabel("Actual Quality")
plt.ylabel('Predicted Quality')
plt.show()

```



```
[130]: import numpy as np
import matplotlib.pyplot as plt
sorted_idx = np.argsort(y_test)
y_test_sorted = np.array(y_test)[sorted_idx]
y_pred_sorted = y_pred[sorted_idx]

# Scatter Plot (Actual vs Predicted)
plt.figure(figsize=(8, 5))
plt.scatter(y_test_sorted, y_pred_sorted, color = 'red', alpha=0.5)
plt.plot(y_test_sorted, y_pred_sorted, color = 'blue')
plt.title('Sorted Actual Sales vs Predicted Sales')
plt.xlabel('Predicted Sales')
plt.ylabel('Actual Sales')
plt.show()
```



Poly_4_Weight-Height

February 6, 2025

```
[166]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import math
```

```
[167]: url = r"Practice-2 Weight-Height Polynomial Dataset.csv"

df = pd.read_csv(url)

df.head()
```

```
[167]:
```

	Weight	Height
0	69.963210	96.644532
1	116.057145	196.156340
2	98.559515	145.862047
3	87.892679	121.157923
4	52.481491	68.971292

1 Data Preprocessing

Checking for null values

```
[168]: df.isna().sum()
```

```
[168]: Weight    0
Height    0
dtype: int64
```

Checking Statistics

```
[169]: df.describe(include='all')
```

```
[169]:
```

	Weight	Height
count	50.000000	50.000000
mean	75.673912	111.473633
std	23.110656	39.493803
min	41.646760	68.971292
25%	54.701360	79.966731
50%	74.883900	98.819101
75%	91.988395	129.709758
max	117.592788	202.663424

Checking Info

```
[170]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 50 entries, 0 to 49  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  -----  -  
0   Weight   50 non-null     float64  
1   Height   50 non-null     float64  
dtypes: float64(2)  
memory usage: 932.0 bytes
```

Checking Shape


```
[171]: df.shape
```

```
[171]: (50, 2)
```

Visualizing the Data

```
[172]: sns.pairplot(df)
```

```
[172]: <seaborn.axisgrid.PairGrid at 0x252f1d9db50>
```

Scaling the data

```
[173]: sca = StandardScaler()
X = df['Weight'].values.reshape(-1,1)
```

Creating X and y

```
[174]: X = df.Weight.values
y = df.Height.values
```

Train Test Split

```
[175]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

Reshaping the data

```
[176]: X_train = X_train.reshape(-1, 1)
X_test = X_test.reshape(-1, 1)
```

```
#y_train=y_train.reshape(-1,1)
#y_test=y_test.reshape(-1,1)
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
```

X_train shape: (40, 1)

X_test shape: (10, 1)

y_train shape: (40,)

y_test shape: (10,)

Model Creation

Simple Linear Regression

```
[177]: linear = LinearRegression()
linear.fit(X_train.reshape(-1,1), y_train)
```

```
[177]: LinearRegression()
```

Predicting using Simple Linear Regression

```
[178]: pred_linear = linear.predict(X_test.reshape(-1,1))
```

Evaluating Simple Linear Regression

```
[179]: mse = mean_squared_error(y_test, pred_linear)
rmse = math.sqrt(mse)
r2 = r2_score(y_test, pred_linear)
mae = mean_absolute_error(y_test, pred_linear)

print('MSE:', mse)
print('RMSE:', rmse)
print('R2:', r2)
print('MAE:', mae)
```

MSE: 117.22192004841368


RMSE: 10.826907224522323

R2: 0.8403417139170934

MAE: 8.882234589447728

Visualizing Using Simple Linear Regression

```
[180]: plt.scatter(X_train, y_train, color = 'blue')
plt.plot(X_test, pred_linear, color = 'lightblue')
plt.title("Height Prediction (Linear)")
plt.xlabel("Height_Actual")
plt.ylabel('Height_Predicted')
plt.show()
```



2 Using Polynomial Regression

```
[181]: Poly = PolynomialFeatures(degree=2)
```

Transforming the data

```
[182]: X_train_poly = Poly.fit_transform(X_train)
X_test_poly = Poly.fit_transform(X_test)
#y_train=y_train.reshape(-1,1)
#y_test=y_test.reshape(-1,1)
```

Training the Model

```
[183]: lin = LinearRegression()

lin.fit(X_train_poly, y_train)
```

```
[183]: LinearRegression()
```

Getting Predictions

```
[184]: y_pred = lin.predict(X_test_poly)
```

Evaluating the Model

```
[185]: r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = math.sqrt(mse)

print(f"R2 Score: {r2}")
print(f"Mean Squared Error: {mse}")
print(f"Mean Absolute Error: {mae}")
print(f"Root Mean Squared Error: {rmse}")
```

R2 Score: 0.9717532753949885


Mean Squared Error: 20.73888787431194

Mean Absolute Error: 3.461429854635168


Root Mean Squared Error: 4.553996911978745

Visualizing the Predictions

```
[186]: sns.regplot(x=y_test, y=y_pred, data=df, scatter_kws={"color": "blue"},
               line_kws={"color": "red"}, order=2, label="Test Data")
plt.xlabel('Actual_weight')
plt.ylabel('Predicted_Weight')
plt.legend()
plt.show()
```



```
[187]: plt.scatter(y_test,y_pred, color='blue',alpha=0.6)
plt.plot(y_test, y_pred, color='red', linewidth=2, label='Polynomial Regression_
↳Line')
plt.xlabel('Actual Height')
plt.ylabel('Predicted Height')
plt.title('Actual Height vs Predicted Height')
plt.legend()
plt.show()
```



```
[188]: import numpy as np
sorted_idx = np.argsort(y_test)
Y_test_sorted = np.array(y_test)[sorted_idx]
y_pred_sorted = y_pred[sorted_idx]

plt.scatter(y_test, y_pred, color='blue')
plt.plot(Y_test_sorted, y_pred_sorted, color='lightblue')
plt.title("Actual vs Predicted Height using Polynomial Regression")
plt.xlabel("Actual Height")
plt.ylabel("Predicted Height")
plt.show()
```

