

Poly_4_Weight-Height

February 6, 2025

```
[166]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import math
```

```
[167]: url = r"Practice-2 Weight-Height Polynomial Dataset.csv"

df = pd.read_csv(url)

df.head()
```

```
[167]:
```

	Weight	Height
0	69.963210	96.644532
1	116.057145	196.156340
2	98.559515	145.862047
3	87.892679	121.157923
4	52.481491	68.971292

1 Data Preprocessing

Checking for null values

```
[168]: df.isna().sum()
```

```
[168]: Weight    0
Height    0
dtype: int64
```

Checking Statistics

```
[169]: df.describe(include='all')
```

```
[169]:
```

	Weight	Height
count	50.000000	50.000000
mean	75.673912	111.473633
std	23.110656	39.493803
min	41.646760	68.971292
25%	54.701360	79.966731
50%	74.883900	98.819101
75%	91.988395	129.709758
max	117.592788	202.663424

Checking Info

```
[170]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 50 entries, 0 to 49  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  -----  -  
0   Weight   50 non-null      float64  
1   Height   50 non-null      float64  
dtypes: float64(2)  
memory usage: 932.0 bytes
```

Checking Shape


```
[171]: df.shape
```

```
[171]: (50, 2)
```

Visualizing the Data

```
[172]: sns.pairplot(df)
```

```
[172]: <seaborn.axisgrid.PairGrid at 0x252f1d9db50>
```



Scaling the data

```
[173]: sca = StandardScaler()
X = df['Weight'].values.reshape(-1,1)
```

Creating X and y

```
[174]: X = df.Weight.values
y = df.Height.values
```

Train Test Split

```
[175]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

Reshaping the data

```
[176]: X_train = X_train.reshape(-1, 1)
X_test = X_test.reshape(-1, 1)
```

```
#y_train=y_train.reshape(-1,1)
#y_test=y_test.reshape(-1,1)
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
```

X_train shape: (40, 1)

X_test shape: (10, 1)

y_train shape: (40,)

y_test shape: (10,)

Model Creation

Simple Linear Regression

```
[177]: linear = LinearRegression()
linear.fit(X_train.reshape(-1,1), y_train)
```

```
[177]: LinearRegression()
```

Predicting using Simple Linear Regression

```
[178]: pred_linear = linear.predict(X_test.reshape(-1,1))
```

Evaluating Simple Linear Regression

```
[179]: mse = mean_squared_error(y_test, pred_linear)
rmse = math.sqrt(mse)
r2 = r2_score(y_test, pred_linear)
mae = mean_absolute_error(y_test, pred_linear)

print('MSE:', mse)
print('RMSE:', rmse)
print('R2:', r2)
print('MAE:', mae)
```

MSE: 117.22192004841368

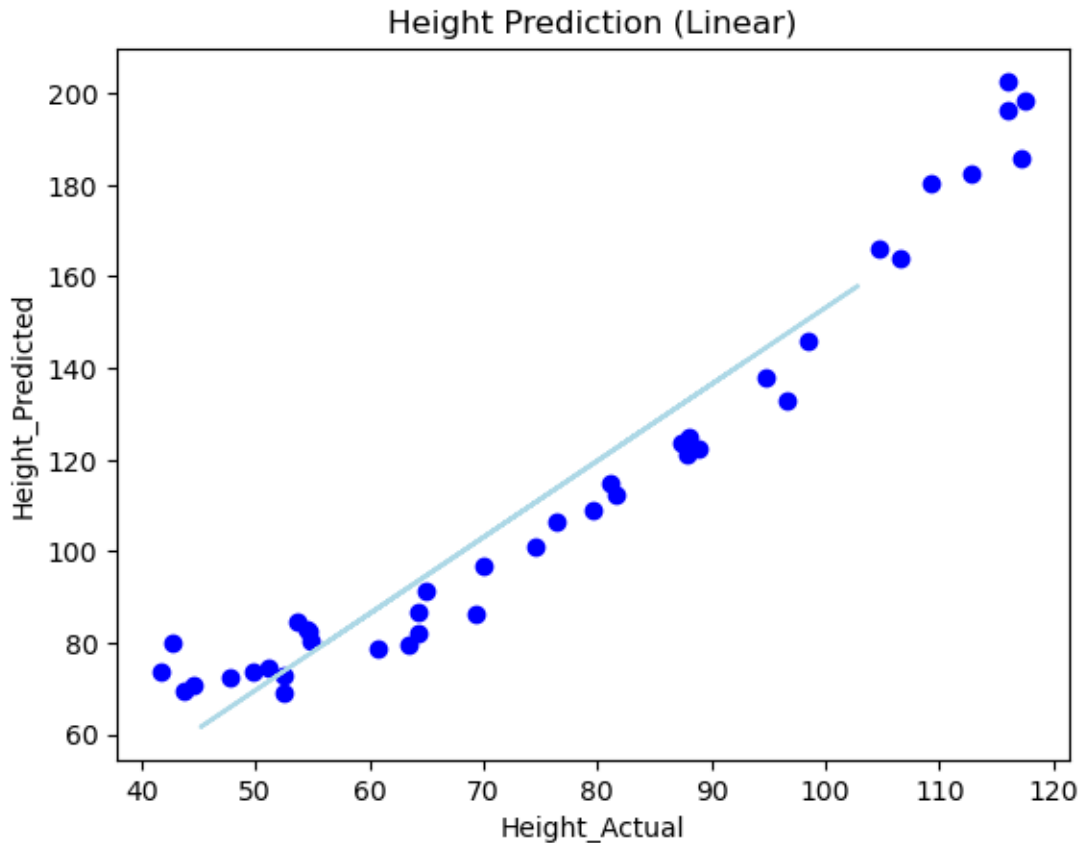
RMSE: 10.826907224522323

R2: 0.8403417139170934

MAE: 8.882234589447728

Visualizing Using Simple Linear Regression

```
[180]: plt.scatter(X_train, y_train, color = 'blue')
plt.plot(X_test, pred_linear, color = 'lightblue')
plt.title("Height Prediction (Linear)")
plt.xlabel("Height_Actual")
plt.ylabel('Height_Predicted')
plt.show()
```



2 Using Polynomial Regression

```
[181]: Poly = PolynomialFeatures(degree=2)
```

Transforming the data

```
[182]: X_train_poly = Poly.fit_transform(X_train)
X_test_poly = Poly.fit_transform(X_test)
#y_train=y_train.reshape(-1,1)
#y_test=y_test.reshape(-1,1)
```

Training the Model

```
[183]: lin = LinearRegression()

lin.fit(X_train_poly, y_train)
```

```
[183]: LinearRegression()
```

Getting Predictions

```
[184]: y_pred = lin.predict(X_test_poly)
```

Evaluating the Model

```
[185]: r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = math.sqrt(mse)

print(f"R2 Score: {r2}")
print(f"Mean Squared Error: {mse}")
print(f"Mean Absolute Error: {mae}")
print(f"Root Mean Squared Error: {rmse}")
```

R2 Score: 0.9717532753949885


Mean Squared Error: 20.73888787431194

Mean Absolute Error: 3.461429854635168

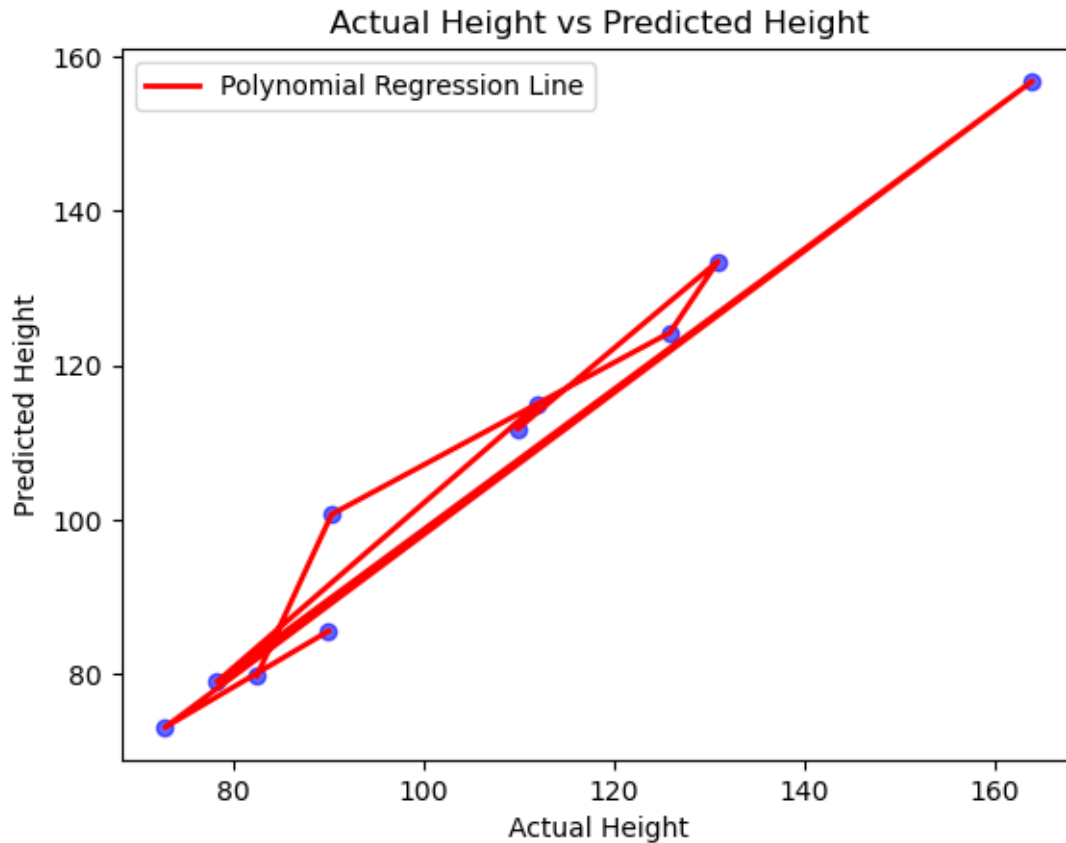
Root Mean Squared Error: 4.553996911978745

Visualizing the Predictions

```
[186]: sns.regplot(x=y_test, y=y_pred, data=df, scatter_kws={"color": "blue"},
               line_kws={"color": "red"}, order=2, label="Test Data")
plt.xlabel('Actual_weight')
plt.ylabel('Predicted_Weight')
plt.legend()
plt.show()
```



```
[187]: plt.scatter(y_test,y_pred, color='blue',alpha=0.6)
plt.plot(y_test, y_pred, color='red', linewidth=2, label='Polynomial Regression_
↳Line')
plt.xlabel('Actual Height')
plt.ylabel('Predicted Height')
plt.title('Actual Height vs Predicted Height')
plt.legend()
plt.show()
```



```
[188]: import numpy as np
sorted_idx = np.argsort(y_test)
Y_test_sorted = np.array(y_test)[sorted_idx]
y_pred_sorted = y_pred[sorted_idx]

plt.scatter(y_test, y_pred, color='blue')
plt.plot(Y_test_sorted, y_pred_sorted, color='lightblue')
plt.title("Actual vs Predicted Height using Polynomial Regression")
plt.xlabel("Actual Height")
plt.ylabel("Predicted Height")
plt.show()
```