

Lasso_Ridge_Class_Assignment

February 16, 2025

1 Ridge And Lasso Regression

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, r2_score, mean_squared_error
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import GridSearchCV
```

```
[2]: df = pd.read_csv('Advertising.csv')
df.head()
```

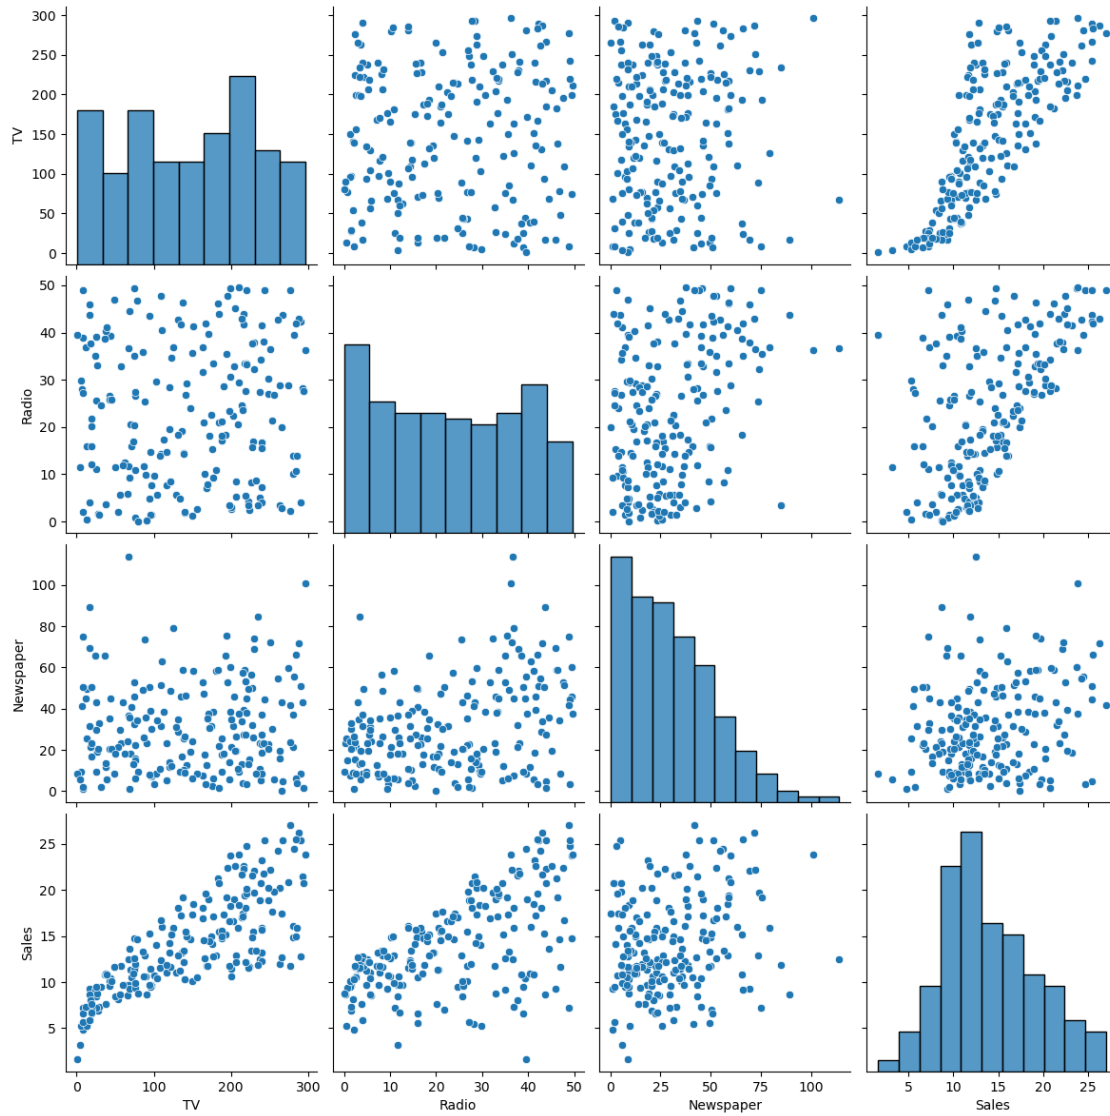
```
[2]: Unnamed: 0    TV  Radio  Newspaper  Sales
0           1  230.1    37.8         69.2   22.1
1           2   44.5    39.3         45.1   10.4
2           3   17.2    45.9         69.3    9.3
3           4  151.5    41.3         58.5   18.5
4           5  180.8    10.8         58.4   12.9
```

```
[3]: df = df.iloc[ : , 1: ]
df.head()
```

```
[3]:      TV  Radio  Newspaper  Sales
0  230.1   37.8         69.2   22.1
1   44.5   39.3         45.1   10.4
2   17.2   45.9         69.3    9.3
3  151.5   41.3         58.5   18.5
4  180.8   10.8         58.4   12.9
```

```
[4]: sns.pairplot(data = df , height=3)
```

```
[4]: <seaborn.axisgrid.PairGrid at 0x2b028cd7350>
```



```
[5]: x = df[['TV', 'Radio', 'Newspaper']]
     y = df['Sales']
```

```
[6]: x.head()
```

```
[6]:
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

```
[7]: y.head()
```

```
[7]: 0    22.1
      1    10.4
      2     9.3
      3    18.5
      4    12.9
      Name: Sales, dtype: float64
```

```
[8]: X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.
      ↪20,random_state=42)
```

```
[9]: from statsmodels.stats.outliers_influence import variance_inflation_factor
      def calc_vif(X):
          vif = pd.DataFrame()
          vif["VIF"] = [variance_inflation_factor(X.values,i) for i in range(X.
          ↪shape[1])]
          return(vif)
```

```
[10]: X = df.iloc[:, :-1]
      calc_vif(X)
```

```
[10]:      VIF
      0  2.486772
      1  3.285462
      2  3.055245
```

```
[11]: alpha_values = np.linspace(-3, 3, 10)
      print(alpha_values)
```

```
[-3.         -2.33333333 -1.66666667 -1.         -0.33333333  0.33333333
  1.         1.66666667  2.33333333  3.         ]
```

```
[12]: from sklearn.linear_model import Ridge
      # Ridge Regression with hyperparameter tuning
      ridge = Ridge()
      # Define the hyperparameter grid for alpha values (regularization strength)
      alpha_values = { 'alpha': np.logspace(-3, 3, 10)} # 10 values from 10^-3 to 10^3

      ridge_cv = GridSearchCV(ridge, alpha_values, cv=5,
          ↪scoring='neg_mean_squared_error')
      ridge_cv.fit(X_train, y_train) # Train Ridge regression model with
          ↪cross-validation

      GridSearchCV(cv=5, estimator=Ridge(), param_grid={'alpha': np.logspace(-3, 3,
          ↪10)}, scoring='neg_mean_squared_error')
```

```
[12]: GridSearchCV(cv=5, estimator=Ridge(),
          param_grid={'alpha': array([1.00000000e-03, 4.64158883e-03,
          2.15443469e-02, 1.00000000e-01,
```

```

4.64158883e-01, 2.15443469e+00, 1.00000000e+01, 4.64158883e+01,
2.15443469e+02, 1.00000000e+03])),
scoring='neg_mean_squared_error')

```

```

[13]: # Best Ridge Model
best_ridge = ridge_cv.best_estimator_
y_pred_ridge = best_ridge.predict(X_test) # Predictions on test data

```

```

[14]: # Ridge RMSE
ridge_rmse = np.sqrt(mean_squared_error(y_test, y_pred_ridge))
r2_ridge = r2_score(y_test, y_pred_ridge)
print("Best ridge Alpha :", ridge_cv.best_params_['alpha'])
print("RMSE :", ridge_rmse)
print("R2 :", r2_ridge)

```

```

Best ridge Alpha : 0.001
RMSE : 1.7815996608176399
R2 : 0.8994380241817195

```

```

[15]: from sklearn.linear_model import Lasso
# Lasso Regression with hyperparameter tuning
lasso = Lasso()
lasso_cv = GridSearchCV(lasso, alpha_values, cv=5,
    ↪scoring='neg_mean_squared_error')
# Define the hyperparameter grid for alpha values (regularization strength)
lasso_cv.fit(X_train, y_train) # Train Lasso regression model with
    ↪cross-validation
GridSearchCV(cv=5, estimator=Lasso(), param_grid={'alpha': np.logspace(-3, -3,
    ↪10)}, scoring='neg_mean_squared_error')

```

```

[15]: GridSearchCV(cv=5, estimator=Lasso(),
    param_grid={'alpha': array([0.001, 0.001, 0.001, 0.001, 0.001,
0.001, 0.001, 0.001, 0.001,
0.001])}),
    scoring='neg_mean_squared_error')

```

```

[16]: # Best Lasso Model
best_lasso = lasso_cv.best_estimator_
y_pred_lasso = best_lasso.predict(X_test) # Predictions on test data

```

```

[17]: # Lasso RMSE
lasso_rmse = np.sqrt(mean_squared_error(y_test, y_pred_lasso))
r2_lasso = r2_score(y_test, y_pred_lasso)
print("Best Lasso Alpha :", lasso_cv.best_params_['alpha'])
print("RMSE :", lasso_rmse)
print("R2 :", r2_lasso)

```

```

Best Lasso Alpha : 2.154434690031882

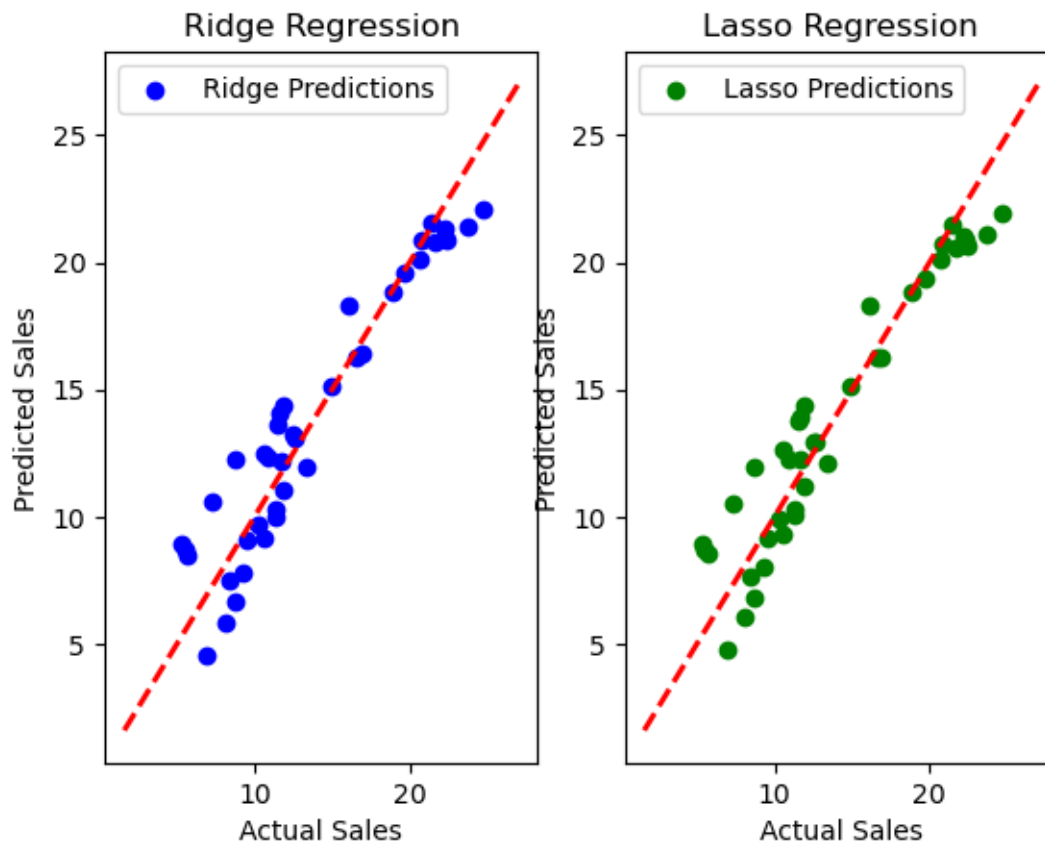
```

RMSE : 1.7677097690307533

R2 : 0.9009999351697155

```
[18]: # Ridge Regression Plot
plt.subplot(1,2,1)
plt.scatter(y_test,y_pred_ridge,color='blue',label="Ridge Predictions")
plt.plot([y.min(),y.max()], [y.min(),y.max()], 'r--',lw=2)
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Ridge Regression')
plt.legend()

# Lasso Regression Plot
plt.subplot(1,2,2)
plt.scatter(y_test,y_pred_lasso,color='green',label="Lasso Predictions")
plt.plot([y.min(),y.max()], [y.min(),y.max()], 'r--',lw=2)
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Lasso Regression')
plt.legend()
plt.show()
```



```
[19]: # Print the coefficients of the lasso with feature names
print("Lasso Coefficients :")
for feature, coef in zip(x.columns, best_lasso.coef_):
    print(f"feature : {feature}, Coefficient , {coef} ")
print("Ridge Coefficients :")
for feature, coef in zip(x.columns, best_lasso.coef_):
    print(f"feature : {feature}, Coefficient , {coef} ")
```

```
Lasso Coefficients :
feature : TV, Coefficient , 0.044516937884577404
feature : Radio, Coefficient , 0.1808414992503233
feature : Newspaper, Coefficient , 0.0
Ridge Coefficients :
feature : TV, Coefficient , 0.044516937884577404
feature : Radio, Coefficient , 0.1808414992503233
feature : Newspaper, Coefficient , 0.0
```