

KNN_Hydropower_Consumption

February 28, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import seaborn as sns
```

```
[2]: df = pd.read_csv("Practice dataset 1 KNN regression Hydropower_Consumption.csv")
```

```
[3]: df
```

```
[3]:
```

	Country	2000	2001	2002	2003	2004	2005	2006	2007	\
0	Afghanistan	312	498	555	63	565	59	637	748	
1	Africa	75246	80864	85181	82873	87405	89066	92241	95341	
2	Albania	4548	3519	3477	5117	5411	5319	4951	276	
3	Algeria	54	69	57	265	251	555	218	226	
4	Angola	903	1007	1132	1229	1733	2197	2638	2472	
..	
148	Uzbekistan	5879	6017	6186	7155	6493	6876	585	6457	
149	Venezuela	62886	60441	59534	60532	70075	77088	81413	83034	
150	Vietnam	14551	1821	18198	0	17818	16535	0	22437	
151	Zambia	7673	7814	8021	8174	8375	8794	9572	9535	
152	Zimbabwe	3227	2968	3786	5305	5466	4866	5257	5329	
	2008	...	2010	2011	2012	2013	2014	2015	2016	\
0	542	...	751	595	71	804	895	989	1025	
1	97157	...	107427	110445	110952	117673	123727	115801	123816	
2	3759	...	7673	4036	4725	6959	4726	5866	7136	
3	283	...	173	378	389	99	193	145	72	
4	3103	...	3666	3967	3734	4719	4991	5037	5757	
..	
148	4386	...	8192	5721	6355	627	6185	602	7327	
149	86713	...	7666	83155	81736	83405	78747	73397	61699	
150	25984	...	28524	41076	53305	5782	62165	57171	66048	
151	9427	...	10331	11368	12227	13148	13902	12907	10915	
152	5651	...	5741	5149	5336	4946	5377	494	2955	

	2017	2018	2019
0	105	105	107
1	130388	132735	0
2	448	448	4018
3	56	117	152
4	7576	7576	8422
..
148	8427	5897	65
149	59296	56987	63267
150	88762	84485	65563
151	12076	12076	11799
152	3929	3929	3592

[153 rows x 21 columns]

```
[4]: df.shape
```

```
[4]: (153, 21)
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 153 entries, 0 to 152
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Country     153 non-null    object
1   2000        153 non-null    int64
2   2001        153 non-null    int64
3   2002        153 non-null    int64
4   2003        153 non-null    int64
5   2004        153 non-null    int64
6   2005        153 non-null    int64
7   2006        153 non-null    int64
8   2007        153 non-null    int64
9   2008        153 non-null    int64
10  2009        153 non-null    int64
11  2010        153 non-null    int64
12  2011        153 non-null    int64
13  2012        153 non-null    int64
14  2013        153 non-null    int64
15  2014        153 non-null    int64
16  2015        153 non-null    int64
17  2016        153 non-null    int64
18  2017        153 non-null    int64
19  2018        153 non-null    int64
20  2019        153 non-null    int64
```

```
dtypes: int64(20), object(1)
memory usage: 25.2+ KB
```

```
[6]: df.isna().sum()
```

```
[6]: Country    0
      2000      0
      2001      0
      2002      0
      2003      0
      2004      0
      2005      0
      2006      0
      2007      0
      2008      0
      2009      0
      2010      0
      2011      0
      2012      0
      2013      0
      2014      0
      2015      0
      2016      0
      2017      0
      2018      0
      2019      0
      dtype: int64
```

```
[7]: df.columns
```

```
[7]: Index(['Country', '2000', '2001', '2002', '2003', '2004', '2005', '2006',
         '2007', '2008', '2009', '2010', '2011', '2012', '2013', '2014', '2015',
         '2016', '2017', '2018', '2019'],
         dtype='object')
```

```
[8]: X = df.drop(columns = ['2019', 'Country'])
      y = df['2019']
```

```
[9]: X = pd.get_dummies(X, drop_first=True)
```

```
[10]: sc = StandardScaler()
      X = sc.fit_transform(X)
```

```
[11]: X_train, X_val, y_train, y_val = train_test_split(X, y, test_size = 0.2,
      ↪ random_state = 42, shuffle = True)
```

```
[12]: params = {
      'n_neighbors': [3,5,7,12],
```

```

    'weights' : ['uniform', 'distance'],
    'metric': ['minkowski', 'manhattan', 'euclidean']
}

```

```
[13]: dia_reg = GridSearchCV(KNeighborsRegressor(), params, cv = 10)
```

```
[14]: dia_reg.fit(X_train, y_train)
```

```

c:\Users\Neil\anaconda3\Lib\site-
packages\joblib\externals\loky\backend\context.py:136: UserWarning: Could not
find the number of physical cores for the following reason:
[WinError 2] The system cannot find the file specified
Returning the number of logical cores instead. You can silence this warning by
setting LOKY_MAX_CPU_COUNT to the number of cores you want to use.
  warnings.warn(
    File "c:\Users\Neil\anaconda3\Lib\site-
packages\joblib\externals\loky\backend\context.py", line 257, in
_count_physical_cores
        cpu_info = subprocess.run(
                    ~~~~~
File "c:\Users\Neil\anaconda3\Lib\subprocess.py", line 548, in run
    with Popen(*popenargs, **kwargs) as process:
        ~~~~~
File "c:\Users\Neil\anaconda3\Lib\subprocess.py", line 1026, in __init__
    self._execute_child(args, executable, preexec_fn, close_fds,
File "c:\Users\Neil\anaconda3\Lib\subprocess.py", line 1538, in _execute_child
    hp, ht, pid, tid = _winapi.CreateProcess(executable, args,
                    ~~~~~

```

```

[14]: GridSearchCV(cv=10, estimator=KNeighborsRegressor(),
                param_grid={'metric': ['minkowski', 'manhattan', 'euclidean'],
                            'n_neighbors': [3, 5, 7, 12],
                            'weights': ['uniform', 'distance']})

```

```
[15]: dia_reg.best_score_
```

```
[15]: -0.7084955694221238
```

```
[16]: dia_reg.best_params_
```

```
[16]: {'metric': 'minkowski', 'n_neighbors': 12, 'weights': 'uniform'}
```

```

[17]: regressor = KNeighborsRegressor(metric = 'manhattan', n_neighbors= 5,
    ↪ weights='distance')
    regressor.fit(X_train, y_train)

```

```
[17]: KNeighborsRegressor(metric='manhattan', weights='distance')
```

```
[18]: y_pred = regressor.predict(X_val)
```

```
[19]: rmse = np.sqrt(mean_squared_error(y_val, y_pred))  
rmse
```

```
[19]: 3884.94001975831
```

```
[20]: from sklearn.metrics import mean_squared_error  
# Calculate and display Mean Squared Error (MSE)  
mse = mean_squared_error(y_val, y_pred)  
print("MSE value : {:.4f}".format(mse))  
from sklearn.metrics import mean_squared_error  
# Calculate and display Root Mean Squared Error (RMSE)  
mse = mean_squared_error(y_val, y_pred)  
rmse = np.sqrt(mse)  
print("RMSE value : {:.4f}".format(rmse))  
  
from sklearn.metrics import r2_score  
# Calculate and display R2 score (R2)  
print("R2 score value : {:.4f}".format(r2_score(y_val, y_pred)))  
  
from sklearn.metrics import mean_absolute_error  
# Calculate and display Mean Absolute Error (MAE)  
mae = mean_absolute_error(y_val, y_pred)  
print("MAE value : {:.4f}".format(mae))
```

```
MSE value : 15092758.9571
```

```
RMSE value : 3884.9400
```

```
R2 score value : 0.8700
```

```
MAE value : 1572.8438
```

```
[21]: plt.figure(figsize=(8,6))  
sns.regplot(x=y_val, y=y_pred, scatter_kws={"s":50}, line_kws={"color": "red"})  
plt.xlabel("Actual Power")  
plt.ylabel("Predicted Power")  
plt.title("KNN Regression: Predicted vs Actual Power")  
plt.grid(True)  
plt.show()
```

